Tokens vs quota consumed

Dots are observed window totals. Lines show the average rate, not a measured per-task conversion.

Muse and Codex: same 22 Oct window. Claude: separate 20 Oct window. Cached input and background reviews included.
High = proportional 5× allowance scenario. Different work/model mixes; throughput does not establish equal-quality task value.